

## CH 13: Logistic Regression

### Motivation

Let's assume that we have access to the underlying candy face off data.

Consider the following model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

where  $y = 1$  if the  $i^{th}$  candy beats smarties and  $y_i = 0$  if  $i^{th}$  candy does not beat smarties,  $x_i$  is an indicator variable that denotes whether the  $i^{th}$  candy has chocolate, and  $\epsilon_i \sim N(0, \sigma^2)$ .

**Q:** What issues might we have with this model?

**Q:** What are some possible solutions?

Logistic regression is a special case of a *generalized linear model*

## Logistic Regression

The logistic function maps an input from the unit range (0,1) to the real line:

$$\text{logit}(x) = \log\left(\frac{x}{1-x}\right)$$

More importantly, the inverse-logit function maps a continuous variable to the unit range (0,1)

$$\text{logit}(x)^{-1} = \frac{\exp(x)}{1 + \exp(x)}$$

The `qlogis` (for logit) and `plogis` (inverse-logit) functions in R can be used for this calculation. For instance `plogis(1) = 0.7310586`.

Formally, the inverse-logistic function is used as part of the GLM:

$$y \sim \text{Bernoulli} \tag{1}$$

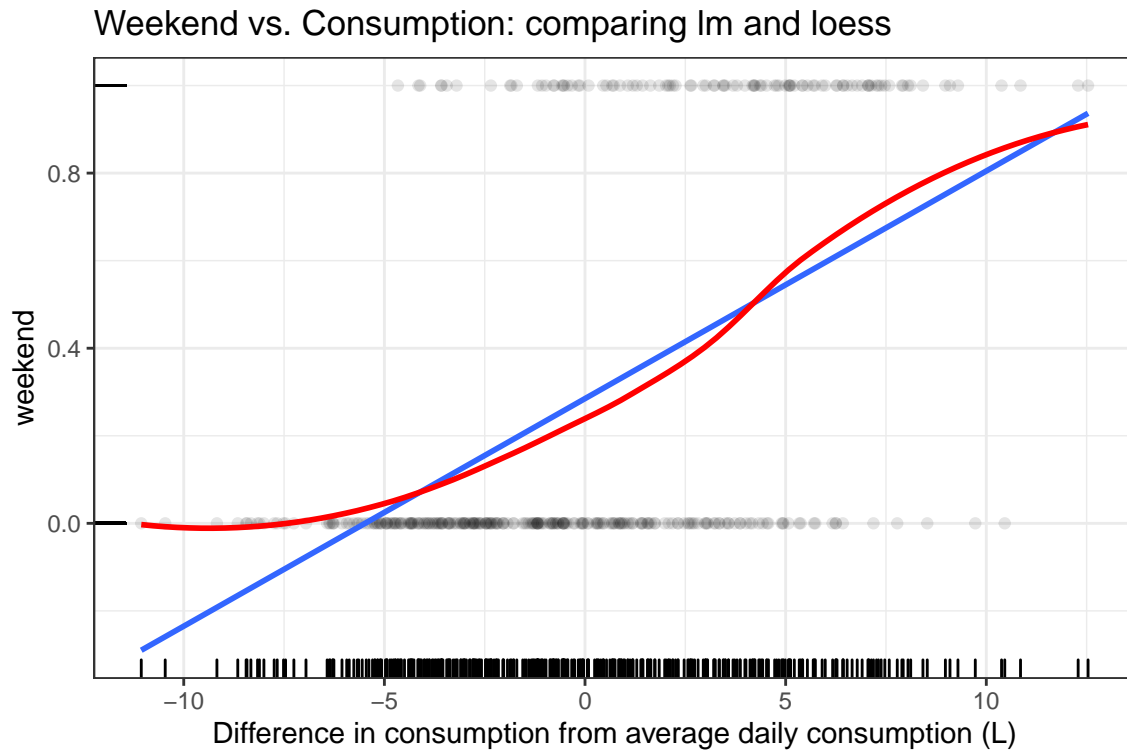
$$\Pr(y_i = 1) = \pi_i = \text{logit}^{-1}(X\underline{\beta}) \tag{2}$$

*Note there is not an  $\epsilon$  term in this model. The randomness comes from the Bernoulli distribution.*

Recall the `beer` dataset, but now instead of trying to model consumption, let's consider whether a day is a weekday or weekend.

```
beer <- read_csv('http://math.montana.edu/ahoegh/Data/Brazil_cerveja.csv') %>% mutate(consumed = consumed - mean(consumed))

beer %>% ggplot(aes(y = weekend, x = consumed)) +
  geom_point(alpha = .1) +
  geom_smooth(formula = 'y~x', method = 'lm', se = F) +
  geom_smooth(formula = 'y~x', method = 'loess', color = 'red', se = F) +
  geom_rug() + ggtitle('Weekend vs. Consumption: comparing lm and loess') +
  theme_bw() + xlab('Difference in consumption from average daily consumption (L)')
```



```
bayes_logistic <- stan_glm(weekend ~ consumed, data = beer,
  family = binomial(link = "logit"), refresh = 0)
```

```
freq_logistic <- glm(weekend ~ consumed, data = beer,
  family = binomial(link = "logit"))
```

Now how to interpret the model coefficients?

```
bayes_logistic
```

```
## stan_glm
## family:      binomial [logit]
## formula:     weekend ~ consumed
## observations: 365
## predictors:  2
## -----
##              Median MAD_SD
## (Intercept) -1.2    0.2
## consumed     0.3    0.0
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

```
summary(freq_logistic)
```

```
##
## Call:
## glm(formula = weekend ~ consumed, family = binomial(link = "logit"),
##      data = beer)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0968  -0.6859  -0.4178   0.7367   2.3624
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.24466    0.15059  -8.265  <2e-16 ***
## consumed     0.31791    0.03773   8.427  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 436.21  on 364  degrees of freedom
## Residual deviance: 333.74  on 363  degrees of freedom
## AIC: 337.74
##
## Number of Fisher Scoring iterations: 5
```

Interpreting the coefficients can be challenging due to the non-linear relationship between the outcome and the predictors.

### Predictive interpretation

One way to interpret the coefficients is in a predictive standpoint. For instance, consider an day with average consumption, then the probability of a weekend would be `invlogit(-1.2 + 0.3 * 0) = 0.23`, where as the probability of a day with 10 more liters of consumption (relative to an average day) would have a weekend probability of `invlogit(-1.2 + 0.3 * 10) = 0.86`

Of course, we should always think about uncertainty, so we can extract simulations from the model.

`posterior_linpred` was useful with regression

```
new_data <- data.frame(consumed = c(0,10))
posterior_sims <- posterior_linpred(bayes_logistic, newdata = new_data)
summary(posterior_sims)
```

```
##           1           2
## Min.      :-1.7585  Min.    :0.917
## 1st Qu.   :-1.3492  1st Qu.:1.728
## Median    :-1.2448  Median :1.950
## Mean      :-1.2482  Mean    :1.960
## 3rd Qu.   :-1.1436  3rd Qu.:2.181
## Max.      :-0.7179  Max.    :3.282
```

*This doesn't return probabilities, so we need to consider `posterior_epred` instead*

```
posterior_sims <- posterior_epred(bayes_logistic, newdata = new_data)
summary(posterior_sims)
```

```
##           1           2
## Min.      :0.1470  Min.    :0.7144
## 1st Qu.   :0.2060  1st Qu.:0.8492
## Median    :0.2236  Median :0.8754
## Mean      :0.2241  Mean    :0.8718
## 3rd Qu.   :0.2417  3rd Qu.:0.8985
## Max.      :0.3279  Max.    :0.9638
```

It can also be useful to consider predictions of an individual data point. *This is how you would conduct posterior predictive checks.*

```
new_obs <- posterior_predict(bayes_logistic, newdata = new_data)
head(new_obs)
```

```
##      1 2
## [1,] 0 0
## [2,] 1 1
## [3,] 1 1
## [4,] 1 0
## [5,] 0 1
## [6,] 0 1
```

```
colMeans(new_obs)
```

```
##      1      2
## 0.22050 0.86425
```

## Model Comparison

We can use cross validation in the same manner a standard linear models.

```
loo(bayes_logistic)
```

```
##
## Computed from 4000 by 365 log-likelihood matrix
##
##      Estimate   SE
## elpd_loo -168.9 10.5
## p_loo      2.0  0.2
## looic     337.8 20.9
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

```
temp_model <- stan_glm(weekend~max_tmp, data = beer, refresh=0)
loo(temp_model)
```

```
##
## Computed from 4000 by 365 log-likelihood matrix
##
##      Estimate   SE
## elpd_loo -230.0  9.1
## p_loo      2.4  0.2
## looic     460.0 18.3
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```